

# PainCare AI Model

Advanced AI-powered endometriosis pain management and prediction system with explainable AI capabilities.

## Overview

PainCare AI is a sophisticated machine learning system designed specifically for endometriosis patients, providing real-time pain predictions, personalized treatment recommendations, and explainable AI insights. The system integrates with Firebase for real-time data synchronization and provides a production-ready REST API for mobile applications.

## Key Capabilities

- **Pain Level Prediction:** ML-powered forecasting using Random Forest algorithms
- **Treatment Recommendations:** Personalized suggestions via K-Means clustering
- **Symptom Pattern Analysis:** Advanced temporal analysis with Gradient Boosting
- **Explainable AI (XAI):** SHAP and LIME-based model explanations
- **Real-time Integration:** Firebase-powered live data synchronization
- **Evidence-based Insights:** Integration with medical research databases
- **Production-ready API:** FastAPI with async endpoints and comprehensive error handling

## Architecture

### System Architecture

```
graph TD
    subgraph Client
        MA[Mobile App]
        REST[REST API]
    end
    subgraph Server
        AI[AI Engine]
        FB[Firebase]
        DB[(Real-time DB)]
    end
    MA --- REST
    REST --- AI
    AI --- FB
    FB --- DB
```

Mobile App REST API AI Engine  
(React N.) (FastAPI) (ML Models)

Firebase  
(Real-time DB)

### Model Pipeline

Raw Data Feature Engineering ML Models XAI Layer API Response

Symptoms 37+ Features 3 Algorithms SHAP/LIME JSON  
Diagnostics Temporal RF, KMeans, Feature Predictions  
User Data Patterns GradBoost Importance Explanations

# Machine Learning Models

## 1. Pain Prediction Model

- **Algorithm:** Random Forest Classifier
- **Features:** 37+ engineered features including temporal patterns, symptom combinations
- **Accuracy:** 85%+ on test data
- **Use Case:** Predicts pain levels 1-7 days ahead

## 2. Treatment Recommendation Engine

- **Algorithm:** K-Means Clustering + Collaborative Filtering
- **Silhouette Score:** 0.25
- **Features:** Treatment history, symptom patterns, user preferences
- **Use Case:** Personalized treatment suggestions

## 3. Symptom Analysis Model

- **Algorithm:** Gradient Boosting Regressor
- **R<sup>2</sup> Score:** 0.63
- **Features:** Temporal symptom data, external factors
- **Use Case:** Pattern recognition and trend analysis

## 4. Explainable AI (XAI) Layer

- **SHAP Values:** Feature importance for individual predictions
- **LIME:** Local model explanations
- **Feature Importance:** Global model insights
- **Use Case:** Transparent AI decision-making

# Features & Data Processing

## Core Features (37+)

### PYTHON:

```
# Symptom Features
- pain_level, sleep_hours, energy_level, mood
- stress_level, physical_activity, medication_taken

# Temporal Features
- day_of_week, hour_of_day, days_since_period
- symptom_trend_3d, symptom_trend_7d

# External Factors
- weather_pressure, temperature, humidity
- medication_effectiveness, treatment_adherence

# Engineered Features
- pain_severity_score, quality_of_life_index
- symptom_pattern_stability, medication_response_rate
```

## Data Augmentation

- **Synthetic Data Generation:** SMOTE for balanced datasets
- **Temporal Augmentation:** Time-series data enhancement
- **Feature Scaling:** StandardScaler for numerical features
- **Categorical Encoding:** LabelEncoder for categorical data

## Quick Start

## Prerequisites

### BASH:

```
Python 3.8+
Firebase Account
8GB RAM (recommended)
```

## Installation

### BASH:

```
# Clone repository
git clone <repository-url>
cd PainCare_Model

# Create virtual environment
python -m venv .venv
source .venv/bin/activate # Windows: .venv\Scripts\activate

# Install dependencies
pip install -r requirements.txt

# Configure environment
cp .env.example .env
# Edit .env with your Firebase credentials
```

## Development Setup

### BASH:

```
# Start development server
python run_server.py

# Or with uvicorn directly
uvicorn src.api.main:app --host 0.0.0.0 --port 8000 --reload

# API will be available at: http://localhost:8000
# Interactive docs: http://localhost:8000/docs
```

## Production Deployment

### 1. Docker Deployment

Create **Dockerfile**:

### DOCKERFILE:

```
FROM python:3.11-slim
```

```
WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

COPY . .

EXPOSE 8000

CMD ["uvicorn", "src.api.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Create `docker-compose.yml`:

## YAML:

```
version: '3.8'
services:
  paincare-ai:
    build: .
    ports:
      - "8000:8000"
    environment:
      - FIREBASE_SERVICE_ACCOUNT_PATH=/app/firebase-service-account.json
      - API_HOST=0.0.0.0
      - API_PORT=8000
    volumes:
      - ./firebase-service-account.json:/app/firebase-service-account.json:ro
    restart: unless-stopped
```

Deploy:

## BASH:

```
docker-compose up -d
```

## 2. Cloud Deployment (AWS/GCP/Azure)

### AWS Elastic Beanstalk

## BASH:

```
# Install EB CLI
pip install awsebcli

# Initialize and deploy
eb init paincare-ai
eb create production
```

```
eb deploy
```

## Google Cloud Run

### BASH:

```
# Build and deploy
gcloud builds submit --tag gcr.io/YOUR_PROJECT/paincare-ai
gcloud run deploy --image gcr.io/YOUR_PROJECT/paincare-ai --platform managed
```

## Azure Container Instances

### BASH:

```
# Create resource group
az group create --name paincare-ai --location eastus

# Deploy container
az container create \
  --resource-group paincare-ai \
  --name paincare-ai-api \
  --image your-registry/paincare-ai:latest \
  --dns-name-label paincare-ai \
  --ports 8000
```

## 3. Kubernetes Deployment

Create `k8s-deployment.yaml`:

### YAML:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: paincare-ai
spec:
  replicas: 3
  selector:
    matchLabels:
      app: paincare-ai
  template:
    metadata:
      labels:
        app: paincare-ai
    spec:
      containers:
```

```
- name: paincare-ai
image: your-registry/paincare-ai:latest
ports:
- containerPort: 8000
env:
- name: API_HOST
value: "0.0.0.0"
resources:
requests:
memory: "512Mi"
cpu: "250m"
limits:
memory: "1Gi"
cpu: "500m"
---
apiVersion: v1
kind: Service
metadata:
name: paincare-ai-service
spec:
selector:
app: paincare-ai
ports:
- protocol: TCP
port: 80
targetPort: 8000
type: LoadBalancer
```

Deploy:

## BASH:

```
kubectl apply -f k8s-deployment.yaml
```

# API Reference

## Endpoints

## Health Check

## HTTP:

```
GET /health
```

## Pain Prediction

### HTTP:

```
POST /predict/pain
Content-Type: application/json
```

```
{
  "symptoms": {
    "pain_level": 6,
    "sleep_hours": 5,
    "stress_level": 8,
    "energy_level": 3,
    "mood": 4
  },
  "include_explanation": true
}
```

## Treatment Recommendations

### HTTP:

```
POST /recommend/treatment
Content-Type: application/json
```

```
{
  "symptoms": {...},
  "medical_history": {...},
  "preferences": {...}
}
```

## Model Status

### HTTP:

```
GET /model/status
```

## XAI Explanations

### HTTP:

```
POST /explain/{prediction_id}
```



## Response Format

### JSON:

```
{
  "success": true,
  "data": {
    "prediction": 6.2,
    "confidence": 0.87,
    "explanation": {
      "top_features": ["stress_level", "sleep_hours"],
      "shap_values": {...}
    }
  },
  "timestamp": "2025-09-02T10:30:00Z",
  "model_version": "1.0.0"
}
```

## Security & Authentication

### API Security

### PYTHON:

```
# JWT Authentication
from fastapi.security import HTTPBearer
security = HTTPBearer()

# Rate Limiting
from slowapi import Limiter
limiter = Limiter(key_func=get_remote_address)

# CORS Configuration
app.add_middleware(
    CORSMiddleware,
    allow_origins=["https://yourapp.com"],
    allow_credentials=True,
    allow_methods=["GET", "POST"],
    allow_headers=["*"],
)
```

## Environment Variables

## ENV:

```
# Firebase
FIREBASE_SERVICE_ACCOUNT_PATH=/path/to/service-account.json

# API Security
SECRET_KEY=your-super-secret-key-here
ALGORITHM=HS256
ACCESS_TOKEN_EXPIRE_MINUTES=30

# Production Settings
DEBUG_MODE=False
LOG_LEVEL=INFO
```

# Monitoring & Observability

## Health Checks

### PYTHON:

```
@app.get("/health")
async def health_check():
    return {
        "status": "healthy",
        "model_loaded": ai_model.is_trained,
        "firebase_connected": firebase_service.is_connected(),
        "timestamp": datetime.now().isoformat()
    }
```

## Logging

### PYTHON:

```
import logging

logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
    handlers=[
        logging.FileHandler('paincare_ai.log'),
        logging.StreamHandler()
    ]
)
```

## Metrics Collection

- **Prometheus:** Custom metrics for model performance
- **Grafana:** Dashboards for monitoring
- **Sentry:** Error tracking and performance monitoring

## Testing

### Unit Tests

#### BASH:

```
# Run all tests
pytest tests/

# Run with coverage
pytest tests/ --cov=src --cov-report=html
```

### Load Testing

#### BASH:

```
# Install locust
pip install locust

# Run load tests
locust -f tests/load_test.py --host http://localhost:8000
```

### Model Validation

#### PYTHON:

```
# Cross-validation
scores = cross_val_score(model, X, y, cv=5)

# Performance metrics
from sklearn.metrics import classification_report
print(classification_report(y_true, y_pred))
```

# Performance Optimization

## Model Optimization

- **Model Pruning:** Remove unnecessary features
- **Quantization:** Reduce model size for deployment
- **Caching:** Redis for frequent predictions
- **Batch Processing:** Handle multiple predictions

## API Optimization

### PYTHON:

```
# Async endpoints
@app.post("/predict/pain")
async def predict_pain(request: PainPredictionRequest):
    result = await ai_model.predict_async(request.symptoms)
    return result

# Response caching
from fastapi_cache import FastAPICache
from fastapi_cache.backends.redis import RedisBackend

@cache(expire=300) # 5 minutes
async def get_cached_prediction():
    return await model.predict(data)
```

# CI/CD Pipeline

## GitHub Actions

### YAML:

```
name: CI/CD Pipeline

on:
  push:
```

```

branches: [ main ]
pull_request:
branches: [ main ]

jobs:
test:
runs-on: ubuntu-latest
steps:
- uses: actions/checkout@v2
- name: Set up Python
uses: actions/setup-python@v2
with:
python-version: 3.11
- name: Install dependencies
run: pip install -r requirements.txt
- name: Run tests
run: pytest tests/

deploy:
needs: test
runs-on: ubuntu-latest
if: github.ref == 'refs/heads/main'
steps:
- name: Deploy to production
run: |
# Your deployment script here
docker build -t paincare-ai .
docker push your-registry/paincare-ai:latest

```

## Model Management

### Model Versioning

#### PYTHON:

```

# Save model with version
joblib.dump(model, f'models/paincare_v{VERSION}.joblib')

# Model registry
class ModelRegistry:
def __init__(self):
self.models = {}

def register_model(self, name: str, version: str, model):
self.models[f"{name}_v{version}"] = {
'model': model,
'timestamp': datetime.now(),
'metrics': self.evaluate_model(model)
}

```

## A/B Testing

### PYTHON:

```
@app.post("/predict/ab_test")
async def ab_test_prediction(request: PredictionRequest):
    # Route 50% to new model, 50% to current
    if hash(request.user_id) % 2 == 0:
        return await new_model.predict(request)
    else:
        return await current_model.predict(request)
```

## Documentation

### API Documentation

- **Swagger UI:** <http://localhost:8000/docs>
- **ReDoc:** <http://localhost:8000/redoc>

### Model Documentation

### PYTHON:

```
# Model cards for transparency
model_card = {
    "model_details": {
        "name": "PainCare Pain Predictor",
        "version": "1.0.0",
        "type": "Random Forest Classifier"
    },
    "intended_use": {
        "primary_uses": "Endometriosis pain prediction",
        "primary_users": "Healthcare providers, patients"
    },
    "metrics": {
        "accuracy": 0.85,
        "precision": 0.83,
        "recall": 0.87
    }
}
```

## Contributing

1. **Fork the repository**
2. **Create feature branch**: ``git checkout -b feature/amazing-feature``
3. **Commit changes**: ``git commit -m 'Add amazing feature'``
4. **Push to branch**: ``git push origin feature/amazing-feature``
5. **Open Pull Request**

## Development Guidelines

- Follow PEP 8 style guide
- Add tests for new features
- Update documentation
- Use type hints
- Add docstrings for all functions

## License

This project is licensed under the MIT License - see the LICENSE file for details.

## Support

### Troubleshooting

- **Model not loading**: Check Firebase credentials
- **Slow predictions**: Enable model caching
- **Memory issues**: Reduce batch size

### Contact

- **Email**: support@paincare.ai
- **Issues**: GitHub Issues
- **Documentation**: Wiki

---

**Built with for endometriosis patients worldwide**